CS288 HW1: Language Modeling

Due: 6 February 2023, 11:00PM PST

Overview

The first homework will focus on language modeling. We'll cover classical n-gram language models, sequence modeling in Pytorch, and tokenization schemes. This homework is intended to be approximately representative of the difficulty of future course homeworks.

Background Reading

Please checkout the following resources before beginning this assignment:

- PyTorch: `https://pytorch.org/tutorials/beginner/deep_learning_60min_blitz.html`
- N-Gram Language Models: `https://web.stanford.edu/~jurafsky/slp3/3.pdf`
- Neural Language Models: `https://web.stanford.edu/~jurafsky/slp3/7.pdf`

Assignment

This homework consists of two iPython notebooks, as well as a written report. You can run and edit the iPython notebooks in any way you please (e.g., run it on your local machine, use an online service such as Google Colab or Kaggle, etc.). Below, you'll find a handful of questions. Please answer these questions in LaTeX and save them to a file named `report.pdf`. Submit this file and all outputs from the notebooks into a single `.zip` file for submission to Gradescope.

1. Getting Started with PyTorch
   (a) Complete the notebook and save the following files:
      - `hw1a.ipynb`
      - `predicted_test_outputs_pooling.txt`
      - `predicted_test_outputs_improved.txt`
   (b) Report: please list at least one correct and one incorrect prediction from your improved network, and give a proposed explanation for why the model might have gotten it wrong. Did the pooling network get these examples right?
2. Training Language Models
   (a) Complete the notebook and save the following files:
      - `hw1b.ipynb`
      - `bigram_predictions.npy`
      - `neural_trigram_predictions.npy`

- `lstm_predictions.npy`

(b) Report: please describe the modifications you made to your LSTM and its corresponding perplexity. Include (1) a concise and precise description of the extension that you tried, (2) a motivation for why you believed this approach might improve your model, (3) a discussion of whether the extension was effective and/or an analysis of the results, and (4) a bottom-line summary of your results comparing validation perplexities of your improvement to the original LSTM. This should involve some combination of tables, learning curves, etc. and be at least half a page in length.

Debugging and Bug Reports

We unfortunately do not have the staffing capacity to help with debugging student code. If you believe you have encountered a legitimate error in one of the notebooks, please post it to the Edstem and we'll do our best to correct the issue as soon as possible.

Submission to Gradescope

Please submit the assignment to Gradescope. The course ID is 500812. The entry code is V5X87D.

When you upload your submission to the Gradescope assignment, you should get immediate feedback that confirms your submission was processed correctly. Be sure to check this, as an incorrectly formatted submission could cause the autograder to fail. For this project, you should be able to see your test set accuracies and a confirmation that all required files were found, but you will not be able to see your score until later. Most assignments will be graded primarily on your test set accuracies and written report.

Note that Gradesope will allow you to submit multiple times before the deadline, and we will use the latest submission for grading. Make sure you have the following files (with correct names and extensions):

- `predicted_test_outputs_pooling.txt`
- `predicted_test_outputs_improved.txt`
- `bigram_predictions.npy`
- `neural_trigram_predictions.npy`
- `lstm_predictions.npy`
- `hw1a.ipynb`
- `hw1b.ipynb`
- `report.pdf`